

✓ Exploratory Data Analysis (EDA) and Data Visualization

Module 4, Lab 2: Understanding Your Data

Exploratory Data Analysis (EDA) is one of the most critical steps in any machine learning project. Before building models, you need to understand your data thoroughly. This lab will teach you how to explore datasets, identify patterns, and create meaningful visualizations.

Learning Objectives

By the end of this lab, you will be able to:

- Load and examine datasets using pandas
- Identify data quality issues (missing values, duplicates, outliers)
- Calculate and interpret summary statistics
- Create effective visualizations using matplotlib and seaborn
- Draw insights from data exploration

Business Problem

We'll analyze a customer dataset to understand purchasing behavior and demographics. This type of analysis helps businesses make data-driven decisions about marketing, product development, and customer segmentation.

✓ Setup and Data Loading

```
# Install required packages
!pip install --upgrade pip
!pip install pandas numpy matplotlib seaborn plotly

Requirement already satisfied: pip in /usr/local/lib/python3.12/dist-packages (25.2)
Requirement already satisfied: pandas in /usr/local/lib/python3.12/dist-packages (2.2.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (2.0.2)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-packages (3.10.0)
Requirement already satisfied: seaborn in /usr/local/lib/python3.12/dist-packages (0.13.2)
Requirement already satisfied: plotly in /usr/local/lib/python3.12/dist-packages (5.24.1)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.12/dist-packages (from pandas) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/dist-packages (from pandas) (2025.2)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.3.3)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (4.60.1)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (1.4.9)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (25.0)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (11.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib) (3.2.5)
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.12/dist-packages (from plotly) (8.5.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.8.2->pandas) (1.17.0)
```

```
# Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')

# Set plotting style
plt.style.use('seaborn-v0_8')
sns.set_palette("husl")
%matplotlib inline

# Set display options
pd.set_option('display.max_columns', None)
pd.set_option('display.width', None)

print("Libraries imported successfully!")

Libraries imported successfully!
```

>Loading the Dataset

We'll create a realistic customer dataset for our analysis.

```
# Create a synthetic customer dataset
np.random.seed(42)
n_customers = 1000

# Generate customer data
customer_data = {
    'customer_id': range(1, n_customers + 1),
    'age': np.random.normal(40, 15, n_customers).astype(int),
    'gender': np.random.choice(['Male', 'Female', 'Other'], n_customers, p=[0.48, 0.50, 0.02]),
    'income': np.random.lognormal(10.5, 0.5, n_customers),
    'education': np.random.choice(['High School', 'Bachelor', 'Master', 'PhD'],
                                  n_customers, p=[0.3, 0.4, 0.25, 0.05]),
    'city_tier': np.random.choice(['Tier 1', 'Tier 2', 'Tier 3'],
                                 n_customers, p=[0.3, 0.4, 0.3]),
    'years_as_customer': np.random.exponential(3, n_customers),
    'total_purchases': np.random.poisson(12, n_customers),
    'avg_order_value': np.random.gamma(2, 50, n_customers),
    'satisfaction_score': np.random.normal(7.5, 1.5, n_customers)
}

# Create DataFrame
df = pd.DataFrame(customer_data)

# Add some realistic constraints
df['age'] = np.clip(df['age'], 18, 80)
df['income'] = np.clip(df['income'], 20000, 200000)
df['years_as_customer'] = np.clip(df['years_as_customer'], 0, 15)
df['satisfaction_score'] = np.clip(df['satisfaction_score'], 1, 10)

# Calculate total spending
df['total_spending'] = df['total_purchases'] * df['avg_order_value']

# Introduce some missing values (realistic scenario)
missing_indices = np.random.choice(df.index, size=50, replace=False)
df.loc[missing_indices, 'satisfaction_score'] = np.nan

# Add some duplicates (data quality issue)
duplicate_rows = df.sample(5).copy()
df = pd.concat([df, duplicate_rows], ignore_index=True)

print(f"Dataset created with {len(df)} customers")
print(f"Dataset shape: {df.shape}")

Dataset created with 1005 customers
Dataset shape: (1005, 11)
```

Step 1: Initial Data Exploration

Let's start by getting familiar with our dataset.

```
# Display first few rows
print("First 5 rows of the dataset:")
display(df.head())

print("\nLast 5 rows of the dataset:")
display(df.tail())
```

First 5 rows of the dataset:

	customer_id	age	gender	income	education	city_tier	years_as_customer	total_purchases	avg_order_value	satisfaction_s
0	1	47	Male	31113.449406	High School	Tier 2	0.471771	12	58.249722	9.07
1	2	37	Male	24932.359005	Bachelor	Tier 2	6.289826	22	39.964016	6.90
2	3	49	Female	42599.051008	Bachelor	Tier 1	2.253532	12	56.137914	5.81
3	4	62	Female	70985.096321	High School	Tier 1	1.889171	13	80.487029	6.99
4	5	36	Male	20000.000000	High School	Tier 1	0.338538	8	84.832080	8.85

Last 5 rows of the dataset:

	customer_id	age	gender	income	education	city_tier	years_as_customer	total_purchases	avg_order_value	satisfaction_s
1000	935	46	Female	96478.813777	Master	Tier 2	3.218394	11	137.763089	9.07
1001	644	18	Male	30728.977611	Bachelor	Tier 3	1.151774	12	92.621642	6.90
1002	810	26	Male	20000.000000	Bachelor	Tier 2	0.227681	11	42.907979	5.81
1003	286	18	Male	67377.429094	Bachelor	Tier 1	1.771158	8	31.961654	6.99
1004	22	36	Female	46833.896860	Bachelor	Tier 2	3.994514	11	42.648150	8.85

```
# Get basic information about the dataset
```

```
print("Dataset Info:")
print(df.info())
```

```
print("\nDataset Shape:")
```

```
print(f"Rows: {df.shape[0]}, Columns: {df.shape[1]}")
```

```
print("\nColumn Names:")
```

```
print(df.columns.tolist())
```

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1005 entries, 0 to 1004
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   customer_id     1005 non-null   int64  
 1   age              1005 non-null   int64  
 2   gender            1005 non-null   object  
 3   income             1005 non-null   float64
 4   education          1005 non-null   object  
 5   city_tier          1005 non-null   object  
 6   years_as_customer 1005 non-null   float64
 7   total_purchases    1005 non-null   int64  
 8   avg_order_value    1005 non-null   float64
 9   satisfaction_score 954 non-null   float64
 10  total_spending     1005 non-null   float64
dtypes: float64(5), int64(3), object(3)
memory usage: 86.5+ KB
None
```

Dataset Shape:

```
Rows: 1005, Columns: 11
```

Column Names:

```
['customer_id', 'age', 'gender', 'income', 'education', 'city_tier', 'years_as_customer', 'total_purchases', 'avg_order_value', 'sa
```

```
# Check data types
print("Data Types:")
print(df.dtypes)
```

```
print("\nNumerical Columns:")
numerical_cols = df.select_dtypes(include=[np.number]).columns.tolist()
print(numerical_cols)
```

```
print("\nCategorical Columns:")
categorical_cols = df.select_dtypes(include=['object']).columns.tolist()
print(categorical_cols)
```

```

Data Types:
customer_id      int64
age              int64
gender          object
income        float64
education      object
city_tier      object
years_as_customer  float64
total_purchases  int64
avg_order_value  float64
satisfaction_score float64
total_spending   float64
dtype: object

Numerical Columns:
['customer_id', 'age', 'income', 'years_as_customer', 'total_purchases', 'avg_order_value', 'satisfaction_score', 'total_spending']

Categorical Columns:
['gender', 'education', 'city_tier']

```

▼ Step 2: Data Quality Assessment

Before analyzing the data, we need to identify and understand data quality issues.

```

# Check for missing values
print("Missing Values:")
missing_data = df.isnull().sum()
missing_percent = (missing_data / len(df)) * 100

missing_df = pd.DataFrame({
    'Missing Count': missing_data,
    'Missing Percentage': missing_percent
})
missing_df = missing_df[missing_df['Missing Count'] > 0].sort_values('Missing Count', ascending=False)
print(missing_df)

Missing Values:
            Missing Count  Missing Percentage
satisfaction_score           51             5.074627

```

```

# Check for duplicate rows
print(f"Total rows: {len(df)}")
print(f"Unique rows: {len(df.drop_duplicates())}")
print(f"Duplicate rows: {len(df) - len(df.drop_duplicates())}")

if len(df) != len(df.drop_duplicates()):
    print("\nDuplicate rows found:")
    duplicates = df[df.duplicated(keep=False)]
    print(duplicates.sort_values('customer_id'))

```

```

Total rows: 1005
Unique rows: 1000
Duplicate rows: 5

Duplicate rows found:
   customer_id  age  gender      income education city_tier \
21          22    36  Female  46833.896860  Bachelor  Tier 2
1004         22    36  Female  46833.896860  Bachelor  Tier 2
1003         286   18  Male   67377.429094  Bachelor  Tier 1
285          286   18  Male   67377.429094  Bachelor  Tier 1
643          644   18  Male   30728.977611  Bachelor  Tier 3
1001         644   18  Male   30728.977611  Bachelor  Tier 3
809          810   26  Male  20000.000000  Bachelor  Tier 2
1002         810   26  Male  20000.000000  Bachelor  Tier 2
934          935   46  Female  96478.813777    Master  Tier 2
1000         935   46  Female  96478.813777    Master  Tier 2

   years_as_customer  total_purchases  avg_order_value  satisfaction_score \
21            3.994514                 11       42.648150        8.604621
1004          3.994514                 11       42.648150        8.604621
1003           1.771158                  8       31.961654        7.293824
285           1.771158                  8       31.961654        7.293824
643           1.151774                 12      92.621642        NaN
1001          1.151774                 12      92.621642        NaN
809           0.227681                 11      42.907979        2.973530
1002          0.227681                 11      42.907979        2.973530
934           3.218394                 11     137.763089        9.208171

```

	3.218394	11	137.763089	9.208171
total_spending				
21	469.129645			
1004	469.129645			
1003	255.693232			
285	255.693232			
643	1111.459705			
1001	1111.459705			
809	471.987767			
1002	471.987767			
934	1515.393977			
1000	1515.393977			

```
# Check for outliers using IQR method
def detect_outliers(df, column):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    outliers = df[(df[column] < lower_bound) | (df[column] > upper_bound)]
    return outliers, lower_bound, upper_bound

print("Outlier Detection (using IQR method):")
for col in ['age', 'income', 'total_spending']:
    outliers, lower, upper = detect_outliers(df, col)
    print(f"\n{col}:")
    print(f"  Normal range: {lower:.2f} to {upper:.2f}")
    print(f"  Number of outliers: {len(outliers)}")
    print(f"  Percentage of outliers: {len(outliers)/len(df)*100:.2f}%")
```

Outlier Detection (using IQR method):

```
age:
  Normal range: 1.50 to 77.50
  Number of outliers: 6
  Percentage of outliers: 0.60%

income:
  Normal range: -10132.92 to 87157.18
  Number of outliers: 36
  Percentage of outliers: 3.58%

total_spending:
  Normal range: -913.79 to 3050.00
  Number of outliers: 43
  Percentage of outliers: 4.28%
```

▼ Step 3: Summary Statistics

Let's calculate and interpret summary statistics for our numerical variables.

```
# Basic summary statistics
print("Summary Statistics for Numerical Variables:")
summary_stats = df.describe()
display(summary_stats.round(2))
```

Summary Statistics for Numerical Variables:

	customer_id	age	income	years_as_customer	total_purchases	avg_order_value	satisfaction_score	total_spending
count	1005.00	1005.00	1005.00	1005.00	1005.00	1005.00	954.00	1005.00
mean	500.69	40.10	41604.00	3.01	12.00	100.63	7.45	1194.66
std	289.10	13.89	20933.79	2.97	3.42	70.04	1.40	886.97
min	1.00	18.00	20000.00	0.01	3.00	1.11	2.97	8.90
25%	251.00	30.00	26350.87	0.82	10.00	50.63	6.52	572.63
50%	501.00	40.00	36443.98	2.08	12.00	84.02	7.44	964.52
75%	751.00	49.00	50673.39	4.27	14.00	133.65	8.44	1563.58
max	1000.00	80.00	174359.45	15.00	23.00	464.25	10.00	6268.79

```
# Additional statistics
print("Additional Statistics:")
additional_stats = pd.DataFrame({
    'Skewness': df[numerical_cols].skew(),
    'Kurtosis': df[numerical_cols].kurtosis(),
    'Variance': df[numerical_cols].var()
})
display(additional_stats.round(3))
```

Additional Statistics:

	Skewness	Kurtosis	Variance	
customer_id	-0.002	-1.202	8.357904e+04	
age	0.328	-0.391	1.928100e+02	
income	1.938	6.351	4.382238e+08	
years_as_customer	1.615	2.666	8.814000e+00	
total_purchases	0.328	-0.035	1.168500e+01	
avg_order_value	1.443	2.818	4.905801e+03	
satisfaction_score	-0.237	-0.138	1.961000e+00	
total_spending	1.606	3.718	7.867163e+05	

```
# Summary for categorical variables
print("Summary for Categorical Variables:")
for col in categorical_cols:
    print(f"\n{col}:")
    value_counts = df[col].value_counts()
    percentages = df[col].value_counts(normalize=True) * 100
    summary = pd.DataFrame({
        'Count': value_counts,
        'Percentage': percentages
    })
    print(summary.round(2))
```

Summary for Categorical Variables:

gender:	Count	Percentage
gender		
Female	519	51.64
Male	465	46.27
Other	21	2.09

education:	Count	Percentage
education		
Bachelor	434	43.18
High School	309	30.75
Master	220	21.89
PhD	42	4.18

city_tier:	Count	Percentage
city_tier		
Tier 2	410	40.80
Tier 3	301	29.95
Tier 1	294	29.25

▼ Step 4: Data Visualization

Now let's create visualizations to better understand our data patterns.

▼ 4.1 Distribution of Numerical Variables

```
# Create histograms for numerical variables
fig, axes = plt.subplots(2, 3, figsize=(18, 12))
axes = axes.ravel()

numerical_vars = ['age', 'income', 'years_as_customer', 'total_purchases', 'avg_order_value', 'total_spending']


for i, var in enumerate(numerical_vars):
```

```

axes[i].hist(df[var].dropna(), bins=30, alpha=0.7, color='skyblue', edgecolor='black')
axes[i].set_title(f'Distribution of {var.replace("_", " ").title()}')
axes[i].set_xlabel(var.replace("_", " ").title())
axes[i].set_ylabel('Frequency')
axes[i].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

```




```

# Box plots to identify outliers
fig, axes = plt.subplots(2, 3, figsize=(18, 12))
axes = axes.ravel()

for i, var in enumerate(numerical_vars):
    axes[i].boxplot(df[var].dropna())
    axes[i].set_title(f'Box Plot of {var.replace("_", " ").title()}')
    axes[i].set_xlabel(var.replace("_", " ").title())
    axes[i].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

```




4.2 Categorical Variable Analysis

```
# Bar plots for categorical variables
fig, axes = plt.subplots(1, 3, figsize=(18, 6))

for i, var in enumerate(categorical_cols):
    value_counts = df[var].value_counts()
    axes[i].bar(value_counts.index, value_counts.values, alpha=0.7)
    axes[i].set_title(f'Distribution of {var.replace("_", " ").title()}'')
    axes[i].set_xlabel(var.replace("_", " ").title())
    axes[i].set_ylabel('Count')
    axes[i].tick_params(axis='x', rotation=45)
    axes[i].grid(True, alpha=0.3)


plt.tight_layout()
plt.show()
```



```
# Pie charts for categorical variables
fig, axes = plt.subplots(1, 3, figsize=(18, 6))

for i, var in enumerate(categorical_cols):
    value_counts = df[var].value_counts()
    axes[i].pie(value_counts.values, labels=value_counts.index, autopct='%1.1f%%', startangle=90)
    axes[i].set_title(f'Proportion of {var.replace("_", " ").title()}')

plt.tight_layout()
plt.show()
```



4.3 Correlation Analysis

```
# Correlation matrix
correlation_matrix = df[numerical_cols].corr()

plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0,
            square=True, linewidths=0.5, fmt='.2f')
plt.title('Correlation Matrix of Numerical Variables')
plt.tight_layout()
plt.show()

# Print strong correlations
print("Strong Correlations (|r| > 0.5):")
for i in range(len(correlation_matrix.columns)):
    for j in range(i+1, len(correlation_matrix.columns)):
        corr_value = correlation_matrix.iloc[i, j]
        if abs(corr_value) > 0.5:
            print(f"{correlation_matrix.columns[i]} vs {correlation_matrix.columns[j]}: {corr_value:.3f}")
```



Strong Correlations ($|r| > 0.5$):
 avg_order_value vs total_spending : 0.885

4.4 Relationship Analysis

```
# Scatter plots for key relationships
fig, axes = plt.subplots(2, 2, figsize=(15, 12))

# Income vs Total Spending
axes[0, 0].scatter(df['income'], df['total_spending'], alpha=0.6)
axes[0, 0].set_xlabel('Income')
axes[0, 0].set_ylabel('Total Spending')
axes[0, 0].set_title('Income vs Total Spending')
axes[0, 0].grid(True, alpha=0.3)


# Age vs Total Purchases
axes[0, 1].scatter(df['age'], df['total_purchases'], alpha=0.6)
axes[0, 1].set_xlabel('Age')
axes[0, 1].set_ylabel('Total Purchases')
axes[0, 1].set_title('Age vs Total Purchases')
axes[0, 1].grid(True, alpha=0.3)

# ...
```

```
# years as customer vs satisfaction score
axes[1, 0].scatter(df['years_as_customer'], df['satisfaction_score'], alpha=0.6)
axes[1, 0].set_xlabel('Years as Customer')
axes[1, 0].set_ylabel('Satisfaction Score')
axes[1, 0].set_title('Years as Customer vs Satisfaction Score')
axes[1, 0].grid(True, alpha=0.3)

# Average Order Value vs Total Purchases
axes[1, 1].scatter(df['avg_order_value'], df['total_purchases'], alpha=0.6)
axes[1, 1].set_xlabel('Average Order Value')
axes[1, 1].set_ylabel('Total Purchases')
axes[1, 1].set_title('Average Order Value vs Total Purchases')
axes[1, 1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```



4.5 Group Analysis

```
# Spending patterns by gender
plt.figure(figsize=(15, 5))

plt.subplot(1, 3, 1)
df.boxplot(column='total_spending', by='gender', ax=plt.gca())
plt.title('Total Spending by Gender')
plt.suptitle('') # Remove default title
```


```

plt.subplot(1, 3, 2)
df.boxplot(column='income', by='education', ax=plt.gca())
plt.title('Income by Education Level')
plt.suptitle('') # Remove default title
plt.xticks(rotation=45)

plt.subplot(1, 3, 3)
df.boxplot(column='satisfaction_score', by='city_tier', ax=plt.gca())
plt.title('Satisfaction Score by City Tier')
plt.suptitle('') # Remove default title

plt.tight_layout()
plt.show()

```



```

# Group statistics
print("Average Total Spending by Gender:")
gender_spending = df.groupby('gender')['total_spending'].agg(['mean', 'median', 'std']).round(2)
print(gender_spending)

print("\nAverage Income by Education Level:")
education_income = df.groupby('education')['income'].agg(['mean', 'median', 'std']).round(2)
print(education_income)

print("\nAverage Satisfaction Score by City Tier:")
city_satisfaction = df.groupby('city_tier')['satisfaction_score'].agg(['mean', 'median', 'std']).round(2)
print(city_satisfaction)

Average Total Spending by Gender:
    mean   median     std
gender
Female  1254.28  1034.91  897.43
Male    1123.97   900.81  868.60
Other   1286.33  1091.70  955.32


Average Income by Education Level:
    mean   median     std
education
Bachelor  41635.96 36521.03 20196.63
High School 40884.57 36054.03 21005.57
Master    41657.57 36738.09 20194.18
PhD      46286.02 42204.91 30066.34

Average Satisfaction Score by City Tier:
    mean   median     std
city_tier
Tier 1    7.41    7.45  1.43
Tier 2    7.41    7.43  1.36
Tier 3    7.56    7.42  1.42

```

▼ Step 5: Advanced Visualizations

```
# Pair plot for key numerical variables
key_vars = ['age', 'income', 'total_spending', 'satisfaction_score']
sns.pairplot(df[key_vars + ['gender']].dropna(), hue='gender', diag_kind='hist')
plt.suptitle('Pair Plot of Key Variables by Gender', y=1.02)
plt.show()
```



```
# Create customer segments based on spending and purchases
df['spending_category'] = pd.cut(df['total_spending'],
                                 bins=[0, df['total_spending'].quantile(0.33),
                                       df['total_spending'].quantile(0.67),
                                       df['total_spending'].max()],
                                 labels=['Low Spender', 'Medium Spender', 'High Spender'])

# Visualize segments
plt.figure(figsize=(15, 5))

plt.subplot(1, 3, 1)
segment_counts = df['spending_category'].value_counts()
plt.pie(segment_counts.values, labels=segment_counts.index, autopct='%1.1f%%')
plt.title('Customer Segments by Spending')
```


```

plt.subplot(1, 3, 2)
sns.boxplot(data=df, x='spending_category', y='satisfaction_score')
plt.title('Satisfaction Score by Spending Category')
plt.xticks(rotation=45)

plt.subplot(1, 3, 3)
sns.boxplot(data=df, x='spending_category', y='years_as_customer')
plt.title('Customer Tenure by Spending Category')
plt.xticks(rotation=45)

plt.tight_layout()
plt.show()

```



Step 6: Key Insights and Findings

Let's summarize our key findings from the EDA.

```

# Calculate key business metrics
print("== KEY BUSINESS INSIGHTS ==")
print(f"\n📊 Dataset Overview:")
print(f"  • Total customers analyzed: {len(df)}")
print(f"  • Data quality: {(len(df) - df.isnull().sum().sum()) / (len(df) * len(df.columns)) * 100:.1f}% complete")

print(f"\n💰 Financial Metrics:")
print(f"  • Average customer income: ${df['income'].mean():,.0f}")
print(f"  • Average total spending: ${df['total_spending'].mean():,.0f}")
print(f"  • Average order value: ${df['avg_order_value'].mean():,.0f}")
print(f"  • Total revenue: ${df['total_spending'].sum():,.0f}")

print(f"\n-demographic Customer Demographics:")
print(f"  • Average age: {df['age'].mean():.1f} years")
print(f"  • Gender distribution: {dict(df['gender'].value_counts())}")
print(f"  • Average customer tenure: {df['years_as_customer'].mean():.1f} years")

print(f"\n😊 Customer Satisfaction:")
print(f"  • Average satisfaction score: {df['satisfaction_score'].mean():.1f}/10")
print(f"  • Highly satisfied customers (>8): {len(df[df['satisfaction_score'] > 8])}/{len(df.dropna(subset=['satisfaction_score'])})")

print(f"\n_SEGMENT Customer Segments:")
segment_stats = df.groupby('spending_category').agg({
    'total_spending': 'mean',
    'satisfaction_score': 'mean',
    'years_as_customer': 'mean'
}).round(2)
for segment in segment_stats.index:
    count = len(df[df['spending_category'] == segment])
    print(f"  • {segment}: {count} customers ({count/len(df)*100:.1f}%)")
    print(f"    - Avg spending: ${segment_stats.loc[segment, 'total_spending']:,.0f}")
    print(f"    - Avg satisfaction: {segment_stats.loc[segment, 'satisfaction_score']:.1f}/10")

== KEY BUSINESS INSIGHTS ==
📊 Dataset Overview:
  • Total customers analyzed: 1,005

```

- Data quality: 7.9% complete

🟡 Financial Metrics:

- Average customer income: \$41,604
- Average total spending: \$1,195
- Average order value: \$101
- Total revenue: \$1,200,632

👤 Customer Demographics:

- Average age: 40.1 years
- Gender distribution: {'Female': np.int64(519), 'Male': np.int64(465), 'Other': np.int64(21)}
- Average customer tenure: 3.0 years

😊 Customer Satisfaction:

- Average satisfaction score: 7.5/10
- Highly satisfied customers (>8): 337/954 (35.3%)

⌚ Customer Segments:

- Low Spender: 332 customers (33.0%)
 - Avg spending: \$413
 - Avg satisfaction: 7.4/10
- Medium Spender: 341 customers (33.9%)
 - Avg spending: \$990
 - Avg satisfaction: 7.6/10
- High Spender: 332 customers (33.0%)
 - Avg spending: \$2,186
 - Avg satisfaction: 7.4/10

✓ Challenge: Your Turn to Explore!



Now it's your turn to practice EDA skills. Complete the following tasks:

✓ Challenge 1: Create a new visualization

Create a visualization that shows the relationship between education level and average order value. What insights can you draw?

```
# Your code here for Challenge 1
# Hint: Try using a bar plot or box plot
#box plot
sns.boxplot(data=df, x='education', y='avg_order_value')
plt.title('Relationship between Education Level and Average Order Value')
plt.xticks(rotation=45)
plt.show()

#bar plot
plt.bar(df['education'], df['avg_order_value'])
plt.xlabel('Education Level')
plt.ylabel('Average Order Value')
plt.title('Relationship between Education Level and Average Order Value')
plt.xticks(rotation=45)
plt.show()
```



Challenge 2: Identify the most valuable customer segment

Based on the data, identify which combination of characteristics (gender, education, city_tier) represents the most valuable customers.

◆ Gemini

```
# Your code here for Challenge 2
# Hint: Use groupby with multiple columns and calculate mean total_spending
pd.groupby(df, ['gender', 'education', 'city_tier'])['total_spending'].mean()
plt.show()
mean = df.groupby(['gender', 'education', 'city_tier'])['total_spending'].mean()
print(mean)
```

gender	education	city_tier	total_spending
Female	Bachelor	Tier 1	1295.133923

		Tier 2	1222.582850
		Tier 3	1366.603689
High School		Tier 1	1308.130514
		Tier 2	1238.485579
		Tier 3	1239.692201
Master		Tier 1	1250.129111
		Tier 2	1103.171466
		Tier 3	1184.338234
PhD		Tier 1	1507.224127
		Tier 2	1200.660915
		Tier 3	1676.044660
Male	Bachelor	Tier 1	1149.356421
		Tier 2	1034.436421
		Tier 3	1275.329094
High School		Tier 1	1008.435573
		Tier 2	1200.992972
		Tier 3	954.961831
Master		Tier 1	1081.668150
		Tier 2	1210.082608
		Tier 3	966.436526
PhD		Tier 1	1453.336066
		Tier 2	1752.415907
		Tier 3	792.699085
Other	Bachelor	Tier 1	826.768606
		Tier 2	1775.046820
		Tier 3	883.395896
High School		Tier 1	1282.631723
		Tier 2	1480.706605
		Tier 3	1976.311883
Master		Tier 2	689.761682
		Tier 3	1159.169427
		PhD	Tier 3 1250.529027

Name: total_spending, dtype: float64